



ESRGAN: Enhanced Super-Resolution Generative Adversarial Networks

Transforming low-resolution images into photo-realistic high-resolution masterpieces through advanced deep learning.

Name: Tanishq Sohal

Roll Number: 22010103071

What is Super-Resolution and Why Do We Need It?

Problem: We usually encounter low-resolution (LR) images due to limitations in hardware (older cameras, sensors), transmission bandwidth, or storage constraints.

Goal: Super-Resolution (SR) aims to reconstruct a high-resolution (HR) image from its LR counterpart.

Importance: Enhancing resolution is crucial for:

- Improving visual quality for viewing (e.g., legacy photos/videos).
- Boosting performance of other computer vision tasks (e.g., object recognition on surveillance footage).
- Applications in medical imaging, satellite imagery, security, and more.

Basic Methods Fall Short: Simple methods like Bicubic interpolation are fast but often produce blurry and overly smooth results, lacking fine details.

The Evolution of Image Super-Resolution:

- Traditional Methods**
Bicubic interpolation produced simple but blurry results
- SRCNN (2016)**
First CNN(Convolucional Neural Network) approach with good PSNR (Peak Signal to Noise Ratio) but over-smooth textures
- SRGAN (2017)**
Revolutionary adversarial training but suffered from artifacts and instability
- ESRGAN (2018)**
State-of-the-art quality with RRDB(Residual in Residual Dense Blocks) architecture and relativistic discriminator



Three Key Innovations:



RRDB Architecture:

Residual-in-Residual Dense Blocks enable superior feature extraction and complex texture combinations



Relativistic Discriminator:

Ensures stable training and prevents model collapse during adversarial learning

(technique where two models are trained to compete against each other to improve the performance of one or both models)




Enhanced Perceptual Loss:

Combines VGG features, adversarial loss, and content loss for human-aligned results


(Perc. loss measures how well each network is succeeding in its opposing goal)(Ad. loss function used in machine learning, particularly in GANs, to train two competing neural networks: a generator and a discriminator)

ESRGAN Architecture Design:




Residual in Residual Dense Block (RRDB)


Residual Block (RB)



Residual in Residual Dense Block (RRDB)



Training Methodology:



Training Methodology:

Data & Preprocessing:

- DIV2K dataset: split images into train, test, validation sets.
- LR via bicubic downsampling (4×)
- Augmentation: random crop, flip, rotation
- Normalize the images for better training convergence.

Loss Functions:

- Peak Signal-to-Noise Ratio (PSNR)
- The Structural Similarity Index (SSIM)
(like PSNR), SSIM is designed to better align with the human visual system's perception of quality

Training Strategy:

- Stage 1: Pre-train generator with L1 loss
- Stage 2: Fine-tune with combined loss

Optimization:

- Adam optimizer(*an optimization algorithm used in machine learning to efficiently train neural networks*)
- Batch size 4 with 128×128 HR patches
- Training: Alternating updates between the generator and discriminator.

Results and Outputs:




Input Image

Size: (296, 168)



Super-Resolution (4x)

Size: (1184, 672)



Advantages of ESRGAN:

- Photo-Realistic Textures: Superior quality with natural-looking detail generation.
- Stable Training: Relativistic discriminator helps stabilize the adversarial training process.
- Reduced Artifacts: Removal of Batch Normalization layers minimizes common artifacts seen in other networks.
- Enhanced network capacity with RRDB blocks (*Residual in Residual Dense Blocks*)
Benefit: Allows for very deep networks that can learn complex mappings from LR to HR, effectively synthesizing fine textures.
- Reasonable Inference Speed: Practical inference times, e.g., ~2-3 seconds for a 512 x 512 image on typical GPUs.
- Flexibility: Open-source implementations allow for easy fine-tuning on specific domains (e.g., faces, medical images)

Challenges & Trade-offs:

Computational Requirements:

Requires ~11GB GPU memory during training with 8+ hour training time. Faster inference but still requires GPUs for reasonable speed.

Fixed Scale Factor:

Limited to 4× upscaling;
separate models needed for 2× or 8× scaling

Metric Trade-offs:

Prioritizes perceptual quality over traditional metrics like PSNR/SSIM, which may be lower than distortion-optimized methods.

Hallucination Risk:

May generate details not present in original;
Can potentially generate plausible but incorrect details, especially on inputs unlike the training data.

Real-World Applications:



Healthcare:

MRI/CT scan enhancement, pathology slide analysis, and telemedicine quality improvement



Entertainment:

Film restoration, gaming texture upscaling, and VR/AR content enhancement



Security:

CCTV footage enhancement, license plate recognition, and forensic analysis



Scientific Research:

Satellite imagery, microscopy, and astronomical image processing



Consumer Tech:

Smartphone camera enhancement and 4K→8K display upscaling

Literature Review: Key Studies:

Study	Innovation	Key Limitation
SRCNN (2016)	First deep learning approach with 3-layer CNN	Over-smooth results, poor texture synthesis
SRGAN (2017)	First GAN-based SR with perceptual loss	Checkerboard artifacts, training instability
EDSR (2017)	Very deep residual network, removed batch norm	Excellent PSNR but lacks texture details
ESRGAN (2018)	RRDB + relativistic discriminator, won PIRM 2018	Fixed 4× scale, computationally intensive
Real-ESRGAN (2021)	High-order degradation modeling for real-world SR	More computationally expensive

ESRGAN: Key Takeaways:

Significant Leap: ESRGAN represented a major advancement in generating photo-realistic details for Single Image Super-Resolution.

Key Innovations: Success driven by the synergistic effect of the RRDB architecture, Relativistic GANs, and an improved perceptual loss function.

Perceptual Focus: Successfully shifted focus from purely PSNR optimization towards results that align better with human visual perception.

Impact: Widely adopted and influential, forming the basis for subsequent state-of-the-art models like Real-ESRGAN, with numerous practical applications.

Trade-offs: While powerful, it involves trade-offs regarding computational cost, fixed scale factors, and potential detail hallucination.

References:

ESRGAN Paper: Wang, X., Yu, K., Wu, S., Gu, J., Liu, Y., Dong, C., ... & Change Loy, C. (2018). ESRGAN: Enhanced Super-Resolution Generative Adversarial Networks. In Proceedings of the European Conference on Computer Vision Workshops (ECCVW).

SRGAN Paper: Ledig, C., Theis, L., Huszár, F., Caballero, J., Cunningham, A., Acosta, A., ... & Shi, W. (2017). Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR).

EDSR Paper: Lim, B., Son, S., Kim, H., Nah, S., & Mu Lee, K. (2017). Enhanced Deep Residual Networks for Single Image Super-Resolution. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW).

Real-ESRGAN Paper: Wang, X., Xie, L., Dong, C., & Shan, Y. (2021). Real-ESRGAN: Training Real-World Blind Super-Resolution with Pure Synthetic Data. In Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops (ICCVW).

Thank you !

